

```
In [1]: from google.colab import drive
drive.mount('/content/drive')
```

Mounted at /content/drive

```
In [2]: from google.colab import files
files.upload() # Upload kaggle.json when prompted
```

Choose Files No file selected

Upload widget is only available when the cell has been executed in the current

browser session. Please rerun this cell to enable.

Saving kaggle.json to kaggle.json

```
Out[2]: {'kaggle.json': b'{"username": "hamidrezaamin", "key": "d15526e38f947f4469926effee3832aa"}'}
```

```
In [3]: import os

# Create the .kaggle directory
os.makedirs("/root/.kaggle", exist_ok=True)

# Rename and move the file to the proper location
!mv "kaggle.json" /root/.kaggle/kaggle.json

# Set file permissions
!chmod 600 /root/.kaggle/kaggle.json
```

```
In [4]: # Install the Kaggle CLI
!pip install -q kaggle

# Download the competition dataset
!kaggle competitions download -c dogs-vs-cats
```

Downloading dogs-vs-cats.zip to /content

95% 775M/812M [00:03<00:00, 146MB/s]

100% 812M/812M [00:03<00:00, 219MB/s]

```
In [5]: # Unzip the main dataset
!unzip dogs-vs-cats.zip -d /content
```

Archive: dogs-vs-cats.zip

inflating: /content/sampleSubmission.csv

inflating: /content/test1.zip

inflating: /content/train.zip

```
In [6]: # Unzip the inner train and test archives
!unzip -q /content/train.zip -d /content/train
!unzip -q /content/test1.zip -d /content/test
```

```
In [7]: import shutil

# Create the original_dataset_dir
original_dataset_dir = '/content/original_dataset_dir'
os.makedirs(original_dataset_dir, exist_ok=True)

# Move the train and test folders
shutil.move('/content/train', os.path.join(original_dataset_dir, 'train'))
shutil.move('/content/test', os.path.join(original_dataset_dir, 'test'))
```

```
Out[7]: '/content/original_dataset_dir/test'
```

```
In [8]: # Fix train folder
nested_train = '/content/original_dataset_dir/train/train'
correct_train = '/content/original_dataset_dir/train'

if os.path.exists(nested_train):
    for filename in os.listdir(nested_train):
        shutil.move(os.path.join(nested_train, filename), correct_train)
    os.rmdir(nested_train) # remove the empty 'train' folder

# Fix test folder
nested_test = '/content/original_dataset_dir/test/test1'
correct_test = '/content/original_dataset_dir/test'

if os.path.exists(nested_test):
    for filename in os.listdir(nested_test):
        shutil.move(os.path.join(nested_test, filename), correct_test)
    os.rmdir(nested_test) # remove the empty 'test1' folder
```

```
In [9]: # check the number of cats and dogs in train folder
train_dir = '/content/original_dataset_dir/train'

cat_images = [f for f in os.listdir(train_dir) if f.startswith('cat')]
dog_images = [f for f in os.listdir(train_dir) if f.startswith('dog')]
```

```
print(f"Number of cat images: {len(cat_images)}")
print(f"Number of dog images: {len(dog_images)}")
```

Number of cat images: 12500

Number of dog images: 12500

```
In [10]: # Create the base_dir
base_dir = '/content/base_dir'
os.makedirs(base_dir, exist_ok=True)
```

```
In [11]: # Define the base directory
base_dir = '/content/base_dir'
os.makedirs(base_dir, exist_ok=True) # Create base_dir if it doesn't exist

# Define the subfolders to create
subfolders = ['train', 'test', 'validation']

# Loop and create each subfolder inside base_dir
for folder in subfolders:
    path = os.path.join(base_dir, folder)
    os.makedirs(path, exist_ok=True)
    print(f"Created: {path}")
```

Created: /content/base_dir/train

Created: /content/base_dir/test

Created: /content/base_dir/validation

```
In [12]: # Define base directory
base_dir = '/content/base_dir'
os.makedirs(base_dir, exist_ok=True)

# Top-level subfolders
top_subfolders = ['train', 'test', 'validation']

# Sub-subfolders to create inside each top-level folder
class_folders = ['cat', 'dog']

# Create the full directory structure
for sub in top_subfolders:
    sub_path = os.path.join(base_dir, sub)
    os.makedirs(sub_path, exist_ok=True)

    for class_folder in class_folders:
        class_path = os.path.join(sub_path, class_folder)
        os.makedirs(class_path, exist_ok=True)
        print(f"Created: {class_path}")
```

Created: /content/base_dir/train/cat

Created: /content/base_dir/train/dog

Created: /content/base_dir/test/cat

Created: /content/base_dir/test/dog

Created: /content/base_dir/validation/cat

Created: /content/base_dir/validation/dog

```
In [13]: # Source and target directories
source_dir = '/content/original_dataset_dir/train'
target_base_dir = '/content/base_dir/train'

# List all files in the source directory
all_filenames = os.listdir(source_dir)

# Filter out cat and dog images
cat_filenames = [f for f in all_filenames if f.startswith('cat')]
dog_filenames = [f for f in all_filenames if f.startswith('dog')]

# Sort and select the first 1000 of each
cat_filenames = sorted(cat_filenames)[:1000]
dog_filenames = sorted(dog_filenames)[:1000]

# Copy cat images
for fname in cat_filenames:
    src = os.path.join(source_dir, fname)
    dst = os.path.join(target_base_dir, 'cat', fname)
    shutil.copyfile(src, dst)

# Copy dog images
for fname in dog_filenames:
    src = os.path.join(source_dir, fname)
    dst = os.path.join(target_base_dir, 'dog', fname)
    shutil.copyfile(src, dst)

print("Copied 1000 cat and 1000 dog images to base_dir/train/")
```

Copied 1000 cat and 1000 dog images to base_dir/train/

```
In [14]: # Define source and destination paths
source_dir = '/content/original_dataset_dir/train'
validation_dir = '/content/base_dir/validation'

# List all files in the source directory
all_filenames = os.listdir(source_dir)

# Filter cat and dog images
cat_filenames = sorted([f for f in all_filenames if f.startswith('cat')])
dog_filenames = sorted([f for f in all_filenames if f.startswith('dog')])

# Select images from index 1001 to 2000 (i.e., next 1000 images)
cat_val_filenames = cat_filenames[1000:2000]
dog_val_filenames = dog_filenames[1000:2000]

# Copy cat validation images
for fname in cat_val_filenames:
    src = os.path.join(source_dir, fname)
    dst = os.path.join(validation_dir, 'cat', fname)
    shutil.copyfile(src, dst)

# Copy dog validation images
for fname in dog_val_filenames:
    src = os.path.join(source_dir, fname)
    dst = os.path.join(validation_dir, 'dog', fname)
    shutil.copyfile(src, dst)

print("Copied 1000 cat and 1000 dog validation images to base_dir/validation/")
```

Copied 1000 cat and 1000 dog validation images to base_dir/validation/

```
In [15]: # Define source and target directories
source_dir = '/content/original_dataset_dir/test'
target_dir = '/content/base_dir/test'

# List and sort image files (to ensure consistent order like 1.jpg, 2.jpg, ...)
all_test_images = sorted(os.listdir(source_dir))[:1000]

# Copy first 1000 images
for fname in all_test_images:
    src = os.path.join(source_dir, fname)
    dst = os.path.join(target_dir, fname)
    shutil.copyfile(src, dst)

print("Copied first 1000 test images to base_dir/test/")
```

Copied first 1000 test images to base_dir/test/

These are just a few of the options available (for more, see the Keras documentation). Let's quickly go over what we just wrote:

rotation_range is a value in degrees (0-180), a range within which to randomly rotate pictures. width_shift and height_shift are ranges (as a fraction of total width or height) within which to randomly translate pictures vertically or horizontally. shear_range is for randomly applying shearing transformations. zoom_range is for randomly zooming inside pictures. horizontal_flip is for randomly flipping half of the images horizontally -- relevant when there are no assumptions of horizontal asymmetry (e.g. real-world pictures). fill_mode is the strategy used for filling in newly created pixels, which can appear after a rotation or a width/height shift. Let's take a look at our augmented images:

```
In [16]: from tensorflow.keras.preprocessing.image import ImageDataGenerator

datagen = ImageDataGenerator(
    rotation_range=40,
    width_shift_range=0.2,
    height_shift_range=0.2,
    shear_range=0.2,
    zoom_range=0.2,
    horizontal_flip=True,
    fill_mode='nearest')
```

```
In [17]: # This is module with image processing utilities

import keras.utils as image
import matplotlib.pyplot as plt

train_cats_dir = '/content/base_dir/train/cat'

fnames = [os.path.join(train_cats_dir, fname) for fname in os.listdir(train_cats_dir)]

# We pick one image to "augment"
img_path = fnames[70]
```

```

# Read the image and resize it
img = image.load_img(img_path, target_size=(150, 150))



# Convert it to a Numpy array with shape (150, 150, 3)
x = image.img_to_array(img)



# Reshape it to (1, 150, 150, 3)
x = x.reshape((1,) + x.shape)

# The .flow() command below generates batches of randomly transformed images.
# It will loop indefinitely, so we need to `break` the loop at some point!
i = 0
for batch in datagen.flow(x, batch_size=1):
    plt.figure(i)
    imgplot = plt.imshow(image.array_to_img(batch[0]))
    i += 1
    if i % 4 == 0:
        break

plt.show()

```





If we train a new network using this data augmentation configuration, our network will never see twice the same input. However, the inputs that it sees are still heavily intercorrelated, since they come from a small number of original images -- we cannot produce new information, we can only remix existing information. As such, this might not be quite enough to completely get rid of overfitting.

```
In [18]: # Import necessary Keras modules for building and training the model
from keras import layers, models, optimizers
from keras.layers import BatchNormalization

# Build a Convolutional Neural Network (CNN) using the Sequential API
model = models.Sequential()

# First convolutional block:
# - 32 filters with size 3x3, ReLU activation
# - Input shape is 150x150 RGB images
model.add(layers.Conv2D(32, (3, 3), activation='relu', input_shape=(150, 150, 3)))

# Apply batch normalization to stabilize and speed up training
model.add(BatchNormalization())
```

```

# Downsample the feature maps using max pooling
model.add(layers.MaxPooling2D((2, 2)))

# Second convolutional block:
# - 64 filters with size 3x3
model.add(layers.Conv2D(64, (3, 3), activation='relu'))
model.add(BatchNormalization())
model.add(layers.MaxPooling2D((2, 2)))

# Third convolutional block:
# - 128 filters with size 3x3
model.add(layers.Conv2D(128, (3, 3), activation='relu'))
model.add(BatchNormalization())
model.add(layers.MaxPooling2D((2, 2)))

# Fourth convolutional block:
# - Another 128 filters with size 3x3
model.add(layers.Conv2D(128, (3, 3), activation='relu'))
model.add(BatchNormalization())
model.add(layers.MaxPooling2D((2, 2)))

# Flatten the 3D feature maps to 1D before feeding into dense layers
model.add(layers.Flatten())

# Add dropout layer to reduce overfitting by randomly turning off 50% of neurons during training
model.add(layers.Dropout(0.5))

# Fully connected dense layer with 512 units and ReLU activation
model.add(layers.Dense(512, activation='relu'))

# Output layer with sigmoid activation for binary classification (cat vs dog)
model.add(layers.Dense(1, activation='sigmoid'))

# Compile the model:
# - Use binary_crossentropy for binary classification
# - Use RMSprop optimizer with a small learning rate
# - Track accuracy as the performance metric
model.compile(
    loss='binary_crossentropy',
    optimizer=optimizers.RMSprop(learning_rate=1e-4),
    metrics=['acc']
)

# Display the architecture summary including the number of parameters per layer
model.summary()

```

/usr/local/lib/python3.11/dist-packages/keras/src/layers/convolutional/base_conv.py:107: UserWarning: Do not pass an `input_shape`/`input_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.

```
super().__init__(activity_regularizer=activity_regularizer, **kwargs)
```

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 148, 148, 32)	896
batch_normalization (BatchNormalization)	(None, 148, 148, 32)	128
max_pooling2d (MaxPooling2D)	(None, 74, 74, 32)	0
conv2d_1 (Conv2D)	(None, 72, 72, 64)	18,496
batch_normalization_1 (BatchNormalization)	(None, 72, 72, 64)	256
max_pooling2d_1 (MaxPooling2D)	(None, 36, 36, 64)	0
conv2d_2 (Conv2D)	(None, 34, 34, 128)	73,856
batch_normalization_2 (BatchNormalization)	(None, 34, 34, 128)	512
max_pooling2d_2 (MaxPooling2D)	(None, 17, 17, 128)	0
conv2d_3 (Conv2D)	(None, 15, 15, 128)	147,584
batch_normalization_3 (BatchNormalization)	(None, 15, 15, 128)	512
max_pooling2d_3 (MaxPooling2D)	(None, 7, 7, 128)	0
flatten (Flatten)	(None, 6272)	0
dropout (Dropout)	(None, 6272)	0
dense (Dense)	(None, 512)	3,211,776
dense_1 (Dense)	(None, 1)	513

Total params: 3,454,529 (13.18 MB)

Trainable params: 3,453,825 (13.18 MB)

Non-trainable params: 704 (2.75 KB)

```
In [19]: train_datagen = ImageDataGenerator(
    rescale=1./255,
    rotation_range=40,
    width_shift_range=0.2,
    height_shift_range=0.2,
    shear_range=0.2,
    zoom_range=0.2,
    horizontal_flip=True,)

# Point to the correct directories
train_dir = '/content/base_dir/train'
validation_dir = '/content/base_dir/validation'

# Note that the validation data should not be augmented!
test_datagen = ImageDataGenerator(rescale=1./255)

train_generator = train_datagen.flow_from_directory(
    # This is the target directory
    train_dir,
    # All images will be resized to 150x150
    target_size=(150, 150),
    batch_size=20,
    # Since we use binary_crossentropy loss, we need binary labels
    class_mode='binary')

validation_generator = test_datagen.flow_from_directory(
    validation_dir,
    target_size=(150, 150),
    batch_size=20,
    class_mode='binary')

# Fit model
history = model.fit(
    train_generator,
    epochs=100,
    validation_data=validation_generator)
```

Found 2000 images belonging to 2 classes.

Found 2000 images belonging to 2 classes.

```

/usr/local/lib/python3.11/dist-packages/keras/src/trainers/data_adapters/py_dataset_adapter.py:121: UserWarning:
Your `PyDataset` class should call `super().__init__(**kwargs)` in its constructor. `**kwargs` can include `work
ers`, `use_multiprocessing`, `max_queue_size`. Do not pass these arguments to `fit()`, as they will be ignored.
  self._warn_if_super_not_called()

```

[illegible]

100/100	19s	187ms/step	- acc: 0.7812	- loss: 0.4872	- val_acc: 0.7835	- val_loss: 0.5684
Epoch 41/100						
100/100	17s	175ms/step	- acc: 0.7597	- loss: 0.5196	- val_acc: 0.8050	- val_loss: 0.4513
Epoch 42/100						
100/100	18s	178ms/step	- acc: 0.7987	- loss: 0.4511	- val_acc: 0.7860	- val_loss: 0.5072
Epoch 43/100						
100/100	19s	187ms/step	- acc: 0.7586	- loss: 0.5186	- val_acc: 0.6995	- val_loss: 0.6580
Epoch 44/100						
100/100	20s	197ms/step	- acc: 0.7979	- loss: 0.4596	- val_acc: 0.7510	- val_loss: 0.5777
Epoch 45/100						
100/100	17s	175ms/step	- acc: 0.7908	- loss: 0.4781	- val_acc: 0.7510	- val_loss: 0.5681
Epoch 46/100						
100/100	20s	173ms/step	- acc: 0.7676	- loss: 0.4733	- val_acc: 0.7770	- val_loss: 0.4937
Epoch 47/100						
100/100	18s	176ms/step	- acc: 0.8128	- loss: 0.4198	- val_acc: 0.7735	- val_loss: 0.4915
Epoch 48/100						
100/100	19s	186ms/step	- acc: 0.7788	- loss: 0.4674	- val_acc: 0.7395	- val_loss: 0.5639
Epoch 49/100						
100/100	17s	174ms/step	- acc: 0.8119	- loss: 0.4091	- val_acc: 0.7975	- val_loss: 0.5562
Epoch 50/100						
100/100	18s	177ms/step	- acc: 0.7856	- loss: 0.4435	- val_acc: 0.6805	- val_loss: 0.6545
Epoch 51/100						
100/100	18s	185ms/step	- acc: 0.7882	- loss: 0.4603	- val_acc: 0.7580	- val_loss: 0.6024
Epoch 52/100						
100/100	19s	174ms/step	- acc: 0.7909	- loss: 0.4404	- val_acc: 0.7970	- val_loss: 0.5073
Epoch 53/100						
100/100	17s	173ms/step	- acc: 0.7810	- loss: 0.4713	- val_acc: 0.7920	- val_loss: 0.4785
Epoch 54/100						
100/100	19s	193ms/step	- acc: 0.8084	- loss: 0.4204	- val_acc: 0.7955	- val_loss: 0.5051
Epoch 55/100						
100/100	19s	195ms/step	- acc: 0.8027	- loss: 0.4287	- val_acc: 0.8160	- val_loss: 0.4573
Epoch 56/100						
100/100	20s	197ms/step	- acc: 0.8088	- loss: 0.4213	- val_acc: 0.7685	- val_loss: 0.5505
Epoch 57/100						
100/100	19s	195ms/step	- acc: 0.8107	- loss: 0.4319	- val_acc: 0.7910	- val_loss: 0.4885
Epoch 58/100						
100/100	20s	197ms/step	- acc: 0.8163	- loss: 0.4178	- val_acc: 0.7940	- val_loss: 0.4660
Epoch 59/100						
100/100	20s	196ms/step	- acc: 0.8116	- loss: 0.4438	- val_acc: 0.7990	- val_loss: 0.4584
Epoch 60/100						
100/100	18s	175ms/step	- acc: 0.8017	- loss: 0.4411	- val_acc: 0.7645	- val_loss: 0.5849
Epoch 61/100						
100/100	23s	196ms/step	- acc: 0.8037	- loss: 0.4239	- val_acc: 0.7885	- val_loss: 0.4984
Epoch 62/100						
100/100	20s	196ms/step	- acc: 0.8224	- loss: 0.4113	- val_acc: 0.8175	- val_loss: 0.4595
Epoch 63/100						
100/100	17s	175ms/step	- acc: 0.8074	- loss: 0.4171	- val_acc: 0.7955	- val_loss: 0.5125
Epoch 64/100						
100/100	18s	177ms/step	- acc: 0.8036	- loss: 0.4411	- val_acc: 0.6675	- val_loss: 0.7350
Epoch 65/100						
100/100	18s	185ms/step	- acc: 0.7993	- loss: 0.4214	- val_acc: 0.7700	- val_loss: 0.5234
Epoch 66/100						
100/100	20s	177ms/step	- acc: 0.8106	- loss: 0.4250	- val_acc: 0.7990	- val_loss: 0.4919
Epoch 67/100						
100/100	18s	178ms/step	- acc: 0.8300	- loss: 0.3844	- val_acc: 0.7785	- val_loss: 0.5046
Epoch 68/100						
100/100	20s	205ms/step	- acc: 0.8053	- loss: 0.4309	- val_acc: 0.8260	- val_loss: 0.4364
Epoch 69/100						
100/100	18s	175ms/step	- acc: 0.8136	- loss: 0.3965	- val_acc: 0.7995	- val_loss: 0.4657
Epoch 70/100						
100/100	20s	199ms/step	- acc: 0.8166	- loss: 0.4005	- val_acc: 0.7300	- val_loss: 0.6247
Epoch 71/100						
100/100	18s	176ms/step	- acc: 0.8212	- loss: 0.4061	- val_acc: 0.7750	- val_loss: 0.5828
Epoch 72/100						
100/100	17s	174ms/step	- acc: 0.8359	- loss: 0.3696	- val_acc: 0.7865	- val_loss: 0.4928
Epoch 73/100						
100/100	20s	199ms/step	- acc: 0.8495	- loss: 0.3670	- val_acc: 0.7720	- val_loss: 0.5151
Epoch 74/100						
100/100	18s	176ms/step	- acc: 0.8312	- loss: 0.3958	- val_acc: 0.8110	- val_loss: 0.4636
Epoch 75/100						
100/100	18s	181ms/step	- acc: 0.8122	- loss: 0.4425	- val_acc: 0.8090	- val_loss: 0.4430
Epoch 76/100						
100/100	19s	187ms/step	- acc: 0.8299	- loss: 0.3966	- val_acc: 0.8280	- val_loss: 0.4248
Epoch 77/100						
100/100	18s	177ms/step	- acc: 0.8288	- loss: 0.3886	- val_acc: 0.8165	- val_loss: 0.4259
Epoch 78/100						
100/100	20s	198ms/step	- acc: 0.8323	- loss: 0.3753	- val_acc: 0.7990	- val_loss: 0.5259
Epoch 79/100						
100/100	19s	185ms/step	- acc: 0.8452	- loss: 0.3560	- val_acc: 0.8165	- val_loss: 0.4522
Epoch 80/100						
100/100	18s	175ms/step	- acc: 0.8329	- loss: 0.3714	- val_acc: 0.8200	- val_loss: 0.4328
Epoch 81/100						
100/100	23s	198ms/step	- acc: 0.8404	- loss: 0.3708	- val_acc: 0.8100	- val_loss: 0.4409

Epoch 82/100
100/100 ————— 18s 176ms/step - acc: 0.8356 - loss: 0.3656 - val_acc: 0.7115 - val_loss: 0.7168
Epoch 83/100
100/100 ————— 18s 176ms/step - acc: 0.8235 - loss: 0.3912 - val_acc: 0.8200 - val_loss: 0.4426
Epoch 84/100
100/100 ————— 18s 185ms/step - acc: 0.8477 - loss: 0.3503 - val_acc: 0.8090 - val_loss: 0.4440
Epoch 85/100
100/100 ————— 17s 174ms/step - acc: 0.8263 - loss: 0.3979 - val_acc: 0.8160 - val_loss: 0.4326
Epoch 86/100
100/100 ————— 21s 182ms/step - acc: 0.8224 - loss: 0.3758 - val_acc: 0.7975 - val_loss: 0.4569
Epoch 87/100
100/100 ————— 22s 198ms/step - acc: 0.8463 - loss: 0.3587 - val_acc: 0.8180 - val_loss: 0.4526
Epoch 88/100
100/100 ————— 20s 197ms/step - acc: 0.8569 - loss: 0.3407 - val_acc: 0.8090 - val_loss: 0.4721
Epoch 89/100
100/100 ————— 18s 178ms/step - acc: 0.8514 - loss: 0.3309 - val_acc: 0.8170 - val_loss: 0.4592
Epoch 90/100
100/100 ————— 17s 173ms/step - acc: 0.8339 - loss: 0.3708 - val_acc: 0.7885 - val_loss: 0.5005
Epoch 91/100
100/100 ————— 19s 190ms/step - acc: 0.8447 - loss: 0.3618 - val_acc: 0.7915 - val_loss: 0.4696
Epoch 92/100
100/100 ————— 19s 176ms/step - acc: 0.8508 - loss: 0.3473 - val_acc: 0.8265 - val_loss: 0.4213
Epoch 93/100
100/100 ————— 20s 197ms/step - acc: 0.8518 - loss: 0.3495 - val_acc: 0.8145 - val_loss: 0.4467
Epoch 94/100
100/100 ————— 18s 175ms/step - acc: 0.8423 - loss: 0.3776 - val_acc: 0.8395 - val_loss: 0.4219
Epoch 95/100
100/100 ————— 17s 172ms/step - acc: 0.8555 - loss: 0.3471 - val_acc: 0.8130 - val_loss: 0.5013
Epoch 96/100
100/100 ————— 20s 197ms/step - acc: 0.8473 - loss: 0.3507 - val_acc: 0.8315 - val_loss: 0.4310
Epoch 97/100
100/100 ————— 17s 175ms/step - acc: 0.8434 - loss: 0.4019 - val_acc: 0.8050 - val_loss: 0.4748
Epoch 98/100
100/100 ————— 17s 174ms/step - acc: 0.8368 - loss: 0.3489 - val_acc: 0.7920 - val_loss: 0.4986
Epoch 99/100
100/100 ————— 20s 197ms/step - acc: 0.8467 - loss: 0.3734 - val_acc: 0.8230 - val_loss: 0.4996
Epoch 100/100
100/100 ————— 18s 176ms/step - acc: 0.8404 - loss: 0.3537 - val_acc: 0.8140 - val_loss: 0.4666

```
In [20]: acc = history.history['acc']
val_acc = history.history['val_acc']
loss = history.history['loss']
val_loss = history.history['val_loss']

epochs = range(len(acc))


plt.plot(epochs, acc, 'bo', label='Training acc')
plt.plot(epochs, val_acc, 'b', label='Validation acc')
plt.title('Training and validation accuracy')
plt.legend()

plt.figure()

plt.plot(epochs, loss, 'bo', label='Training loss')
plt.plot(epochs, val_loss, 'b', label='Validation loss')
plt.title('Training and validation loss')
plt.legend()

plt.show()
```

Training and validation accuracy



Training and validation loss

